

## Microsoft VibeVoice - Generate Multi-Speaker Long Podcast with AI Locally



Fahd Mirza  
45.5K subscribers

Join

Subscribe

143



Share

Ask

3,173 views Aug 25, 2025 #vibevoice

This video locally installs Microsoft VibeVoice TTS model for generating expressive, long-form, multi-speaker conversational audio.

https://huggingface.co/microsoft/VibeVoice-1.5B 170%

**Hugging Face** Search models, datasets, users... Models Datasets Spaces Docs Pricing

microsoft/**VibeVoice-1.5B** like 120 Follow Microsoft 14.3k

Text-to-Speech Safetensors English Chinese vibevoice Podcast arxiv:2412.08635 License: mit

Model card Files xet Community 2

Edit model card

### VibeVoice: A Frontier Open-Source Text-to-Speech Model

VibeVoice is a novel framework designed for generating expressive, long-form, multi-speaker conversational audio, such as podcasts, from text. It addresses significant challenges in traditional Text-to-Speech (TTS) systems, particularly in scalability, speaker consistency, and natural turn-taking.

A core innovation of VibeVoice is its use of continuous speech

Downloads last month

Safetensors

Model size 2.7B params Tensor type BF16

Files info

Inference Providers NEW

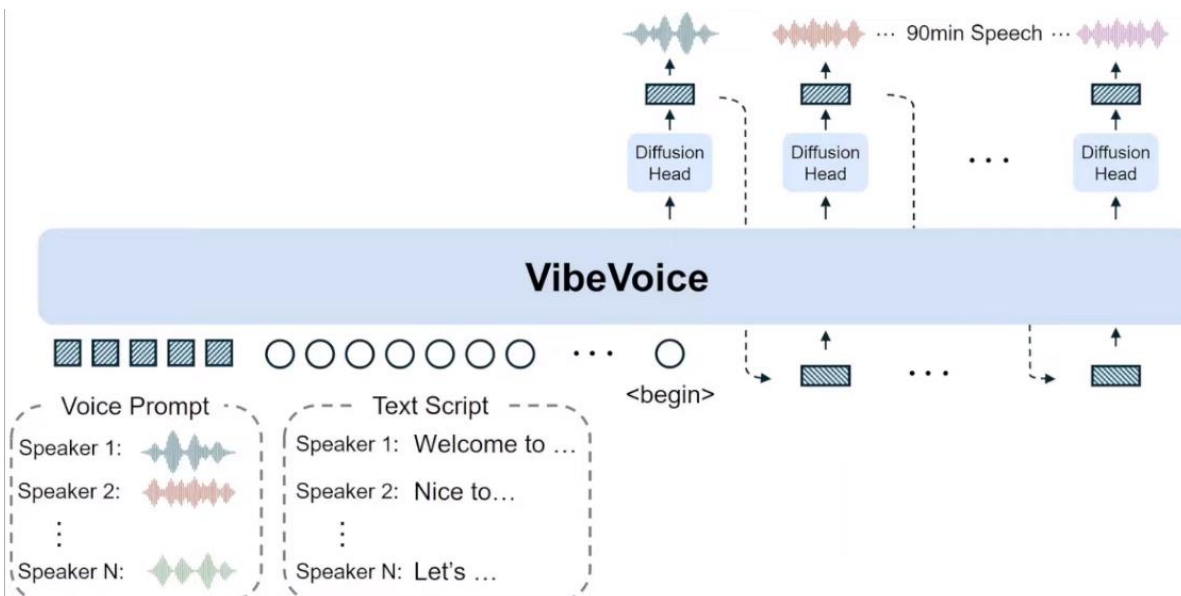
```
Ubuntu@0127-dsm2-ty6k-prxm70113:~$ conda create -n ai python=3.11 -y && conda activate ai
Retrieving notices: ...working... done
Channels:
 - conda-forge
 - defaults
Platform: linux-64
Collecting package metadata (repodata.json): -
```

```
Ubuntu@0127-dsm2-ty6k-prxm70113:~$ conda create -n ai python=3.11 -y && conda activate ai
Retrieving notices: ...working... done
Channels:
 - conda-forge
 - defaults
Platform: linux-64
Collecting package metadata (repodata.json): -
Downloading and Extracting Packages:
Preparing transaction: done
Verifying transaction: done
Executing transaction: done
#
# To activate this environment, use
#
#   $ conda activate ai
#
# To deactivate an active environment, use
#
#   $ conda deactivate
(ai) Ubuntu@0127-dsm2-ty6k-prxm70113:~$
```

```
Ubuntu@0127-dsm2-ty6k-prxm70113:~$ docker --version
Docker version 28.1.1, build 4eba377
(ai) Ubuntu@0127-dsm2-ty6k-prxm70113:~$
```

```
Ubuntu@0127-dsm2-ty6k-prmx70113:~$ sudo chmod 666 /var/run/docker.sock
(ai) Ubuntu@0127-dsm2-ty6k-prmx70113:~$
```

```
Ubuntu@0127-dsm2-ty6k-prmx70113:~$ sudo docker run --privileged --net=host --ipc=host --ulimit mem
lock=1:1 --ulimit stack=1:1 --gpus all --rm -it nvcr.io/nvidia/pytorch:24.07-py3
Unable to find image 'nvcr.io/nvidia/pytorch:24.07-py3' locally
24.07-py3: Pulling from nvidia/pytorch
3713021b0277: Pulling fs layer
34203065b696: Pulling fs layer
4f4fb700ef54: Pulling fs layer
a26410b333db: Waiting
27db04d6bd35: Waiting
aa00ff708333: Waiting
9ec380a307de: Waiting
bc282122a834: Pulling fs layer
ced5e3cb2229: Waiting
de78b163ddcd: Waiting
f132a30b7f59: Waiting
212c0fc745fc: Waiting
6433b0370e40: Waiting
b19cd6da8470: Waiting
32d7c61a0261: Waiting
ea30cbbab0c1: Pulling fs layer
b0d7273c356b: Waiting
aea52c550ecc: Waiting
b6becee36584: Pulling fs layer
e8d08626ff57: Waiting
```



```
root@0127-dsm2-ty6k-prmx70113: /workspace#
Container image Copyright (c) 2024, NVIDIA CORPORATION & AFFILIATES. All rights reserved.
Copyright (c) 2014-2024 Facebook Inc.
Copyright (c) 2011-2014 Idiap Research Institute (Ronan Collobert)
Copyright (c) 2012-2014 Deepmind Technologies (Koray Kavukcuoglu)
Copyright (c) 2011-2012 NEC Laboratories America (Koray Kavukcuoglu)
Copyright (c) 2011-2013 NYU (Clement Farabet)
Copyright (c) 2006-2010 NEC Laboratories America (Ronan Collobert, Leon Bottou, Iain Melvin, Jason Weston)
Copyright (c) 2006 Idiap Research Institute (Samy Bengio)
Copyright (c) 2001-2004 Idiap Research Institute (Ronan Collobert, Samy Bengio, Johnny Mariethoz)
Copyright (c) 2015 Google Inc.
Copyright (c) 2015 Yangqing Jia
Copyright (c) 2013-2016 The Caffe contributors
All rights reserved.

Various files include modifications (c) NVIDIA CORPORATION & AFFILIATES. All rights reserved.

This container image and its contents are governed by the NVIDIA Deep Learning Container License.
By pulling and using the container, you accept the terms and conditions of this license:
https://developer.nvidia.com/ngc/nvidia-deep-learning-container-license

NOTE: CUDA Forward Compatibility mode ENABLED.
Using CUDA 12.5 driver version 555.42.06 with kernel driver version 550.54.14.
See https://docs.nvidia.com/deploy/cuda-compatibility/ for details.

root@0127-dsm2-ty6k-prmx70113: /workspace#
```

```
root@0127-dsm2-ty6k-prmx70113: /workspace/VibeVoice
root@0127-dsm2-ty6k-prmx70113:/workspace# git clone https://github.com/microsoft/VibeVoice.git && cd VibeVoice
Cloning into 'VibeVoice'...
remote: Enumerating objects: 188, done.
remote: Counting objects: 100% (39/39), done.
remote: Compressing objects: 100% (26/26), done.
remote: Total 188 (delta 23), reused 25 (delta 13), pack-reused 149 (from 1)
Receiving objects: 100% (188/188), 85.40 MiB | 52.75 MiB/s, done.
Resolving deltas: 100% (61/61), done.
root@0127-dsm2-ty6k-prmx70113:/workspace/VibeVoice# pip install -e .
Looking in indexes: https://pypi.org/simple, https://pypi.ngc.nvidia.com
Obtaining file:///workspace/VibeVoice
Installing build dependencies ... -
```

```
Building editable for vibevoice (pyproject.toml) ... done
Created wheel for vibevoice: filename=vibevoice-0.0.1-0.editable-py3-none-any.whl size=6237 sha256=7eb68b7000516510aec0838006ffa298813887c184c8a92ee17d9b021eeb4ab1
Stored in directory: /tmp/pip-ephem-wheel-cache-be7cidhx/wheels/ed/13/c5/94d40781d2f1ff3856ce177740c145803f37ce76c858ad72b3
Successfully built vibevoice
Installing collected packages: pydub, ifaddr, brotli, websockets, tomlkit, sniffio, semantic-version, safetensors, ruff, python-multipart, pyee, orjson, ml-collections, hf-xet, h11, groovy, google-crc32c, ffmpeg, dnspython, av, aiofiles, uvicorn, pylibsrtp, huggingface-hub, httpcore, cryptography, anyio, aioice, tokenizers, starlette, pyopenssl, httpx, diffusers, accelerate, transformers, safehttpx, gradio-client, fastapi, aiortc, gradio, vibevoice
Successfully installed accelerate-1.6.0 aiofiles-24.1.0 aioice-0.10.1 aiortc-1.13.0 anyio-4.10.0 av-14.4.0 brotli-1.1.0 cryptography-45.0.6 diffusers-0.35.1 dnspython-2.7.0 fastapi-0.116.1 ffmpeg-0.6.1 google-crc32c-1.7.1 gradio-5.43.1 gradio-client-1.12.1 groovy-0.1.2 h11-0.16.0 hf-xet-1.1.8 httpcore-1.0.9 httpx-0.28.1 huggingface-hub-0.34.4 ifaddr-0.2.0 ml-collections-1.1.0 orjson-3.11.2 pydub-0.25.1 pyee-13.0.0 pylibsrtp-0.12.0 pyopenssl-25.1.0 python-multipart-0.0.20 ruff-0.12.10 safehttpx-0.1.6 safetensors-0.6.2 semantic-version-2.10.0 sniffio-1.3.1 starlette-0.47.3 tokenizers-0.21.4 tomlkit-0.13.3 transformers-4.51.3 uvicorn-0.35.0 vibevoice-0.0.1 websockets-15.0.1
WARNING: Running pip as the 'root' user can result in broken permissions and conflicting behaviour with the system package manager, possibly rendering your system unusable. It is recommended to use a virtual environment instead: https://pip.pypa.io/warnings/venv. Use the --root-user-action option if you know what you are doing and want to suppress this warning.

[notice] A new release of pip is available: 24.1.2 -> 25.2
[notice] To update, run: python -m pip install --upgrade pip
root@0127-dsm2-ty6k-prmx70113:/workspace/VibeVoice#
```

```
root@0127-dsm2-ty6k-prmx70113: /workspace/VibeVoice
root@0127-dsm2-ty6k-prmx70113:/workspace/VibeVoice# python demo/gradio demo.py --model path microsoft/VibeVoice-1.5B --share
```

```
root@0127-dsm2-ty6k-prmx70113: /workspace/VibeVoice
{
  "vae_dim": 64,
  "weight_init_value": 0.01
},
"acoustic vae_dim": 64,
"architectures": [
  "VibeVoiceForConditionalGeneration"
],
"decoder config": {
  "attention dropout": 0.0,
  "hidden_act": "silu",
  "hidden_size": 1536,
  "initializer_range": 0.02,
  "intermediate_size": 8960,
  "max_position_embeddings": 65536,
  "max_window_layers": 28,
  "model type": "qwen2",
  "num_attention_heads": 12,
  "num_hidden_layers": 28,
  "num_key_value_heads": 2,
  "rms_norm_eps": 1e-06,
  "rope_scaling": null,
  "rope_theta": 1000000.0,
  "sliding window": null,
  "tie_word_embeddings": true,
  "torch_dtype": "bfloat16",
  "use_cache": true,

```



```
root@0127-dsm2-tydk-prxm70113: /workspace/VibeVoice
2
},
"fix_std": 0,
"layer_scale_init_value": 1e-06,
"layernorm": "RMSNorm",
"layernorm_elementwise_affine": true,
"layernorm_eps": 1e-05,
"mixer_layer": "depthwise_conv",
"model_type": "vibevoice_semantic_tokenizer",
"pad_mode": "constant",
"std_dist_type": "none",
"vae_dim": 128,
"weight_init_value": 0.01
},
"semantic_vae_dim": 128,
"torch_dtype": "bfloat16",
"transformers_version": "4.51.3"
}

model.safetensors.index.json: 123kB [00:00, 280MB/s]
loading weights file model.safetensors from cache at /root/.cache/huggingface/hub/models--microsoft--VibeVoice-1.5B/snapshots/7a399f30c8b426a6f69bdf378cf8ea725c71f52c/model.safetensors.index.json
Fetching 3 files: 0% | 0/3 [00:00<?, ?it/s]
model-00001-of-00003.safetensors: 14% | 269M/1.98G [00:25<02:24, 11.8MB/s]
model-00002-of-00003.safetensors: 10% | 201M/1.98G [00:25<03:16, 9.05MB/s]
model-00003-of-00003.safetensors: 37% | 537M/1.45G [00:26<00:38, 23.5MB/s]
```

```
root@0127-dsm2-tydk-prxm70113: /workspace/VibeVoice
If your task is similar to the task the model of the checkpoint was trained on, you can already use VibeVoiceForConditionalGenerationInference for predictions without further training.
Generation config file not found, using a generation config created from the model config.
Language model attention: flash attention 2
Found 9 voice files in /workspace/VibeVoice/demo/voices
Available voices: en-Alice_woman, en-Alice_woman_bgm, en-Carter_man, en-Frank_man, en-Maya_woman, in-Samuel_man, zh-Anchen_man_bgm, zh-Bowen_man, zh-Xinran_woman
Loaded example: 1p_Ch2EN.txt with 1 speakers
Loaded example: 1p_abs.txt with 1 speakers
Loaded example: 2p_goat.txt with 2 speakers
Loaded example: 2p_music.txt with 2 speakers
Loaded example: 3p_gpt5.txt with 3 speakers
Skipping 4p_climate_100min.txt: duration 100 minutes exceeds 15-minute limit
Skipping 4p_climate_45min.txt: duration 45 minutes exceeds 15-minute limit
Successfully loaded 5 example scripts
🔥 Launching demo on port 7860
📁 Model path: microsoft/VibeVoice-1.5B
🎤 Available voices: 9
🔴 Streaming mode: ENABLED
🔒 Session isolation: ENABLED
* Running on local URL: http://0.0.0.0:7860
* Running on public URL: https://41b2bef978ce11fdc3.gradio.live

This share link expires in 1 week. For free permanent hosting and GPU upgrades, run `gradio deploy` from the terminal in the working directory to deploy to Hugging Face Spaces (https://huggingface.co/spaces)
```

